Module 3 & 4

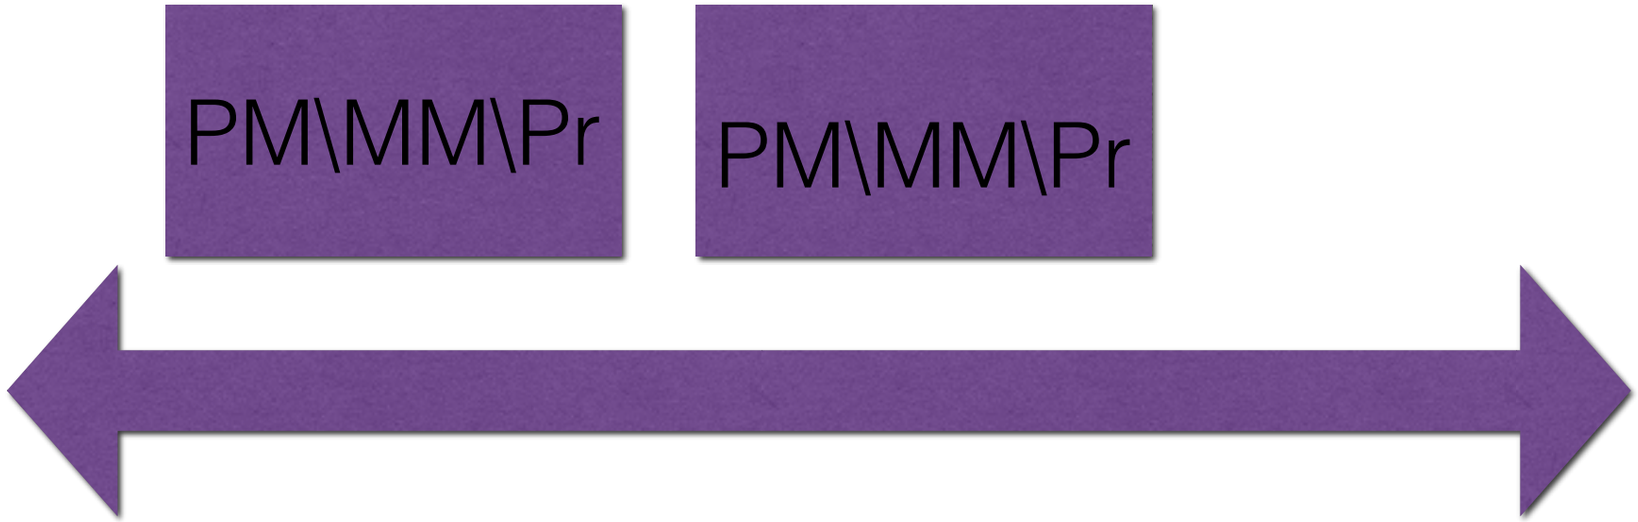
PM\MM\Pr
PM\MM\Pr

JOB1 - Get Data
JOB2 - Process & generate o/p
MAP\REDUCE

4 PM - 5 AM BATCH

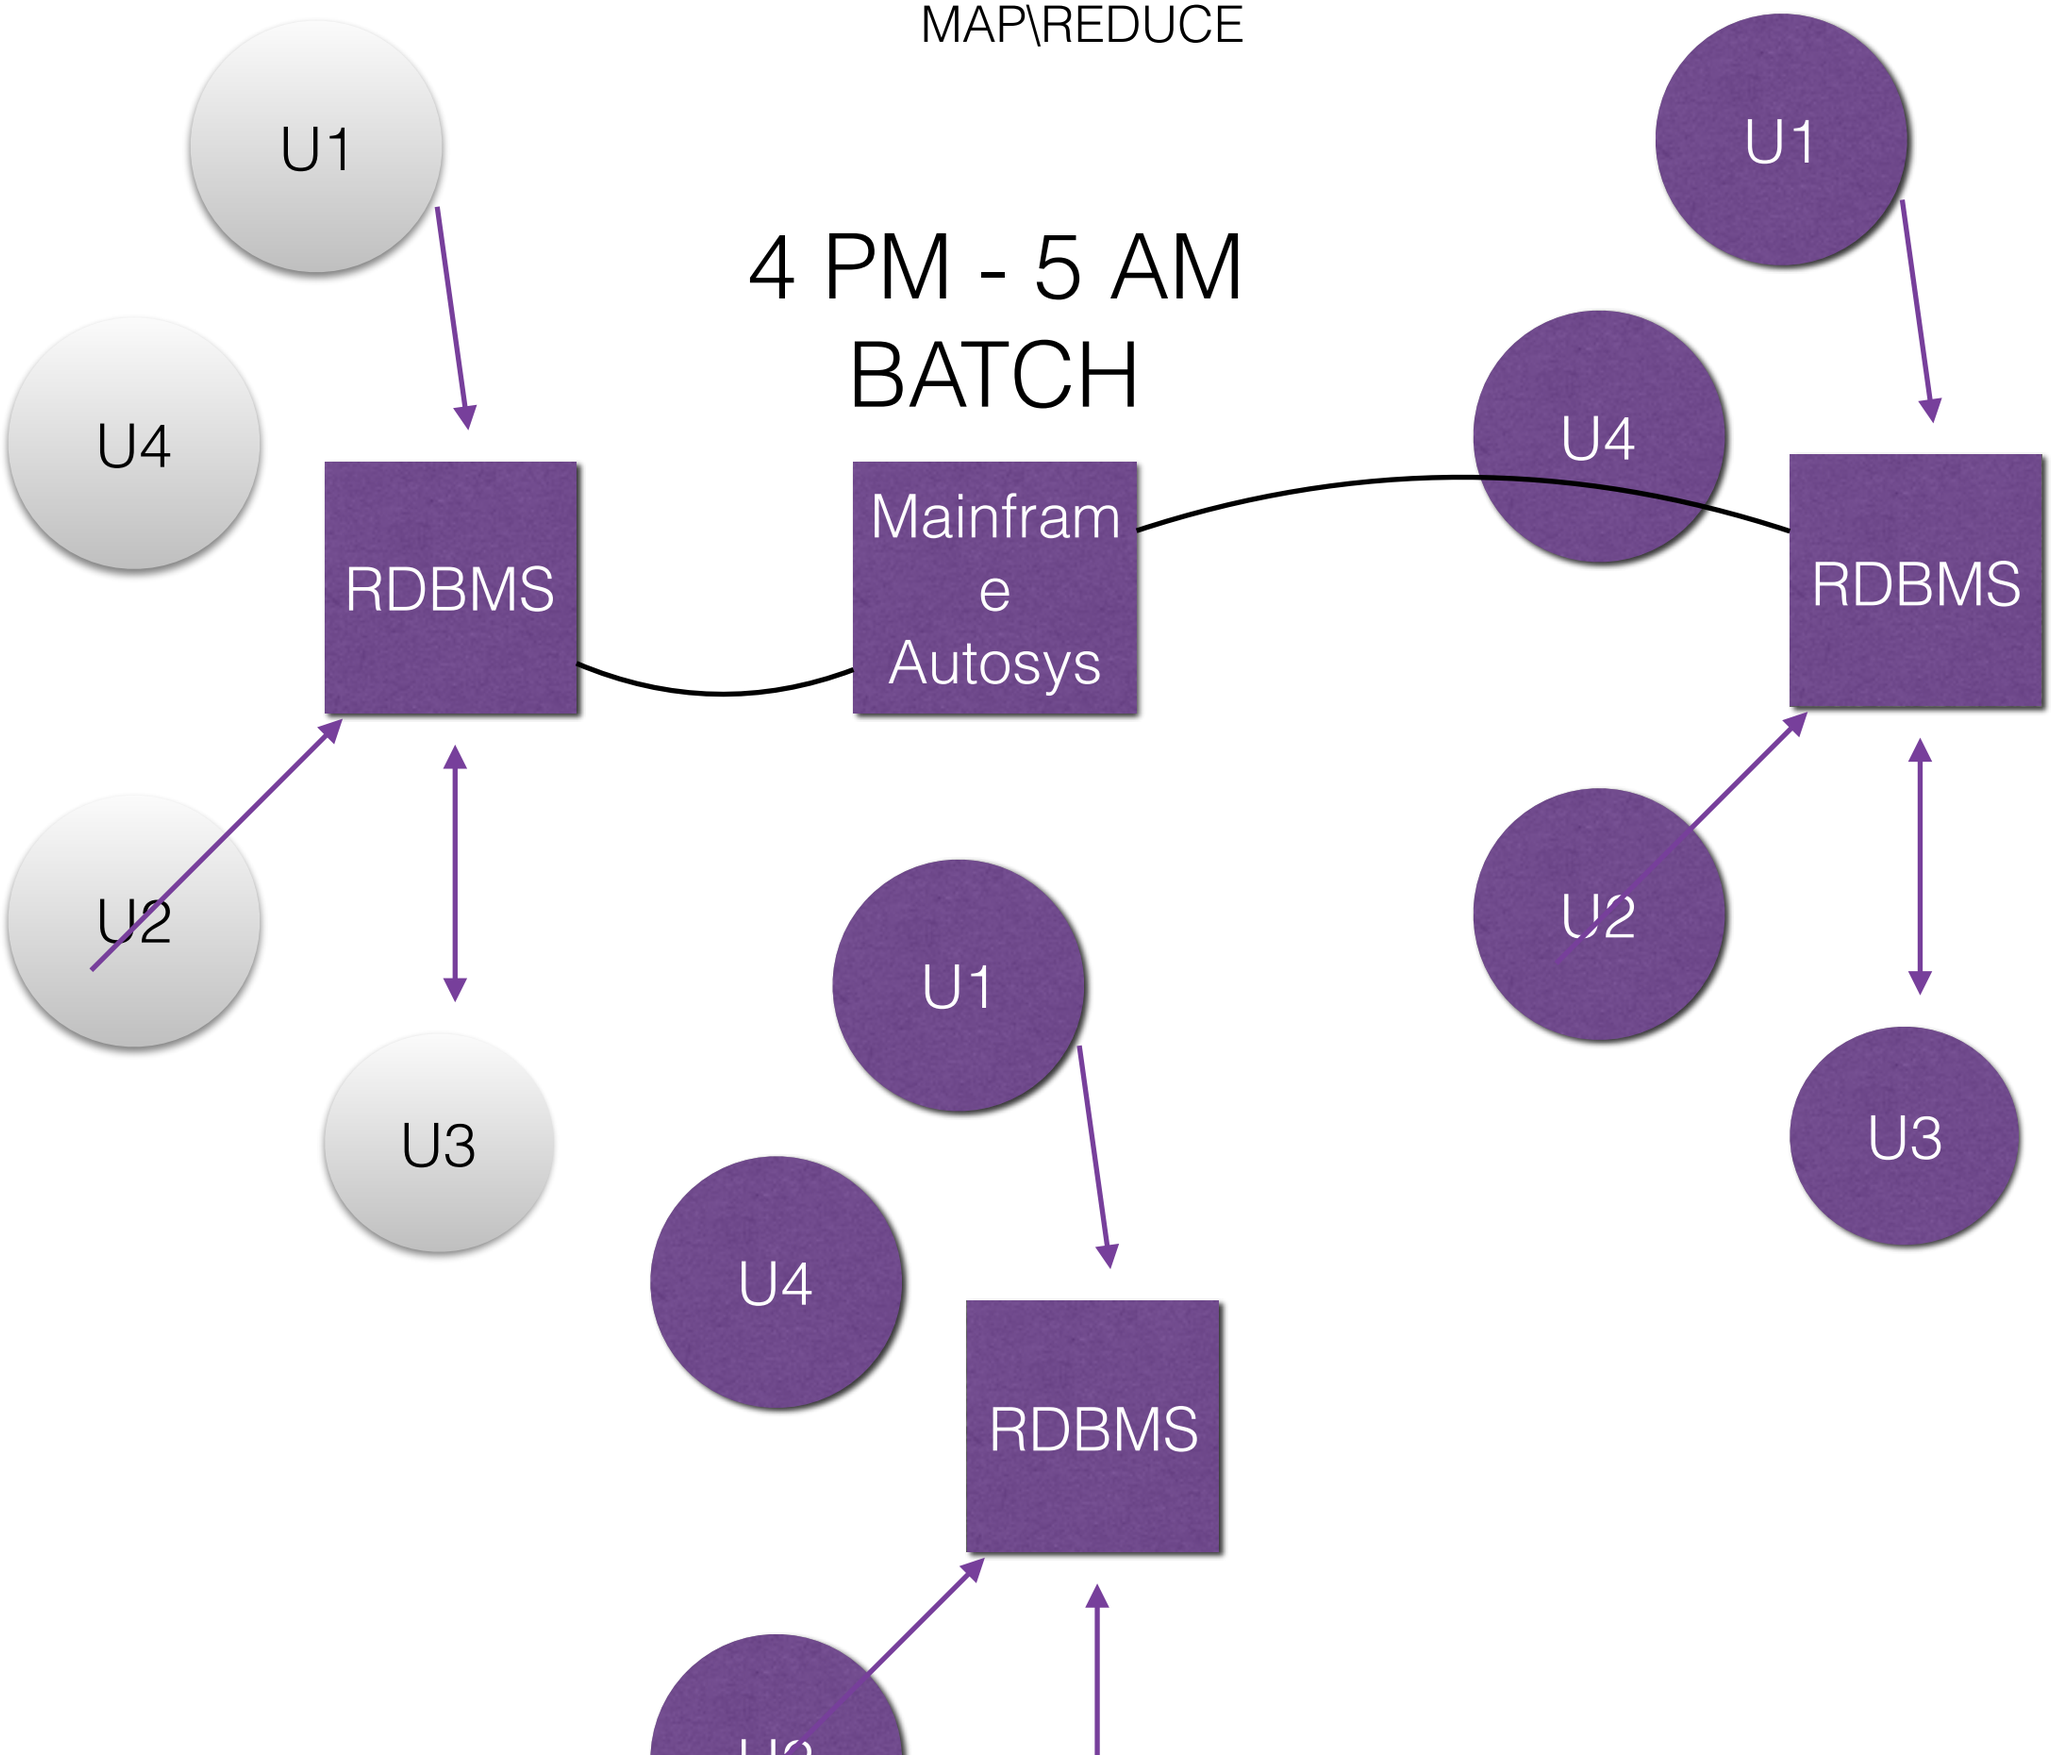

What is Spark?

- Apache Spark™ is a fast and general engine for large-scale data processing.
- It guarantees:

A. ***SPEED:*** Run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk.

- Apache Spark has an advanced DAG execution engine that supports cyclic data flow and in-memory computing.

B. ***EASE OF USE****: Write applications quickly in Java, Scala, Python, R. Spark offers over* ***80 high-level operators*** *that make it easy to build parallel apps. And you can use it interactively from the Scala, Python and R shells.*

C. ***GENERALITY:*** *Combine SQL, streaming, and complex analytics. Spark powers a stack of libraries including* ***SQL and DataFrames****,* ***MLlib for machine learning****,* ***GraphX****, and* ***Spark Streaming.*** *You can combine these libraries seamlessly in the same application.*

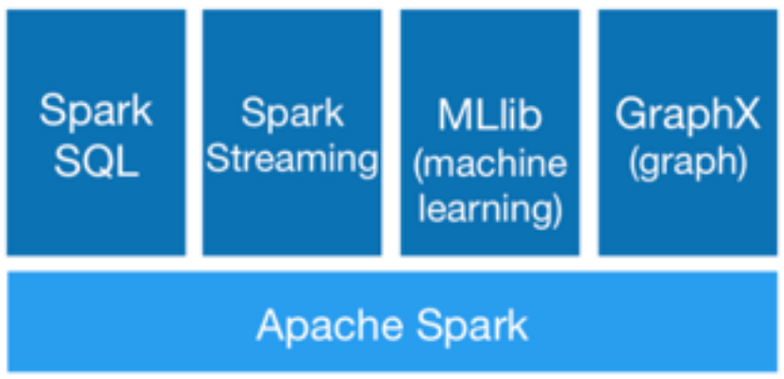

D. ***RUNS EVERYWHERE:*** *Spark runs on* ***Hadoop****,* ***Mesos, standalone****, or in the* ***cloud****. It can access* ***diverse data sources including HDFS, Cassandra, HBase, and S3****.*

SPARK with Scala

SCALA doc for available SPARK Libraries

- https://spark.apache.org/docs/preview/api/scala/index.html#package

Map Reduce Architecture

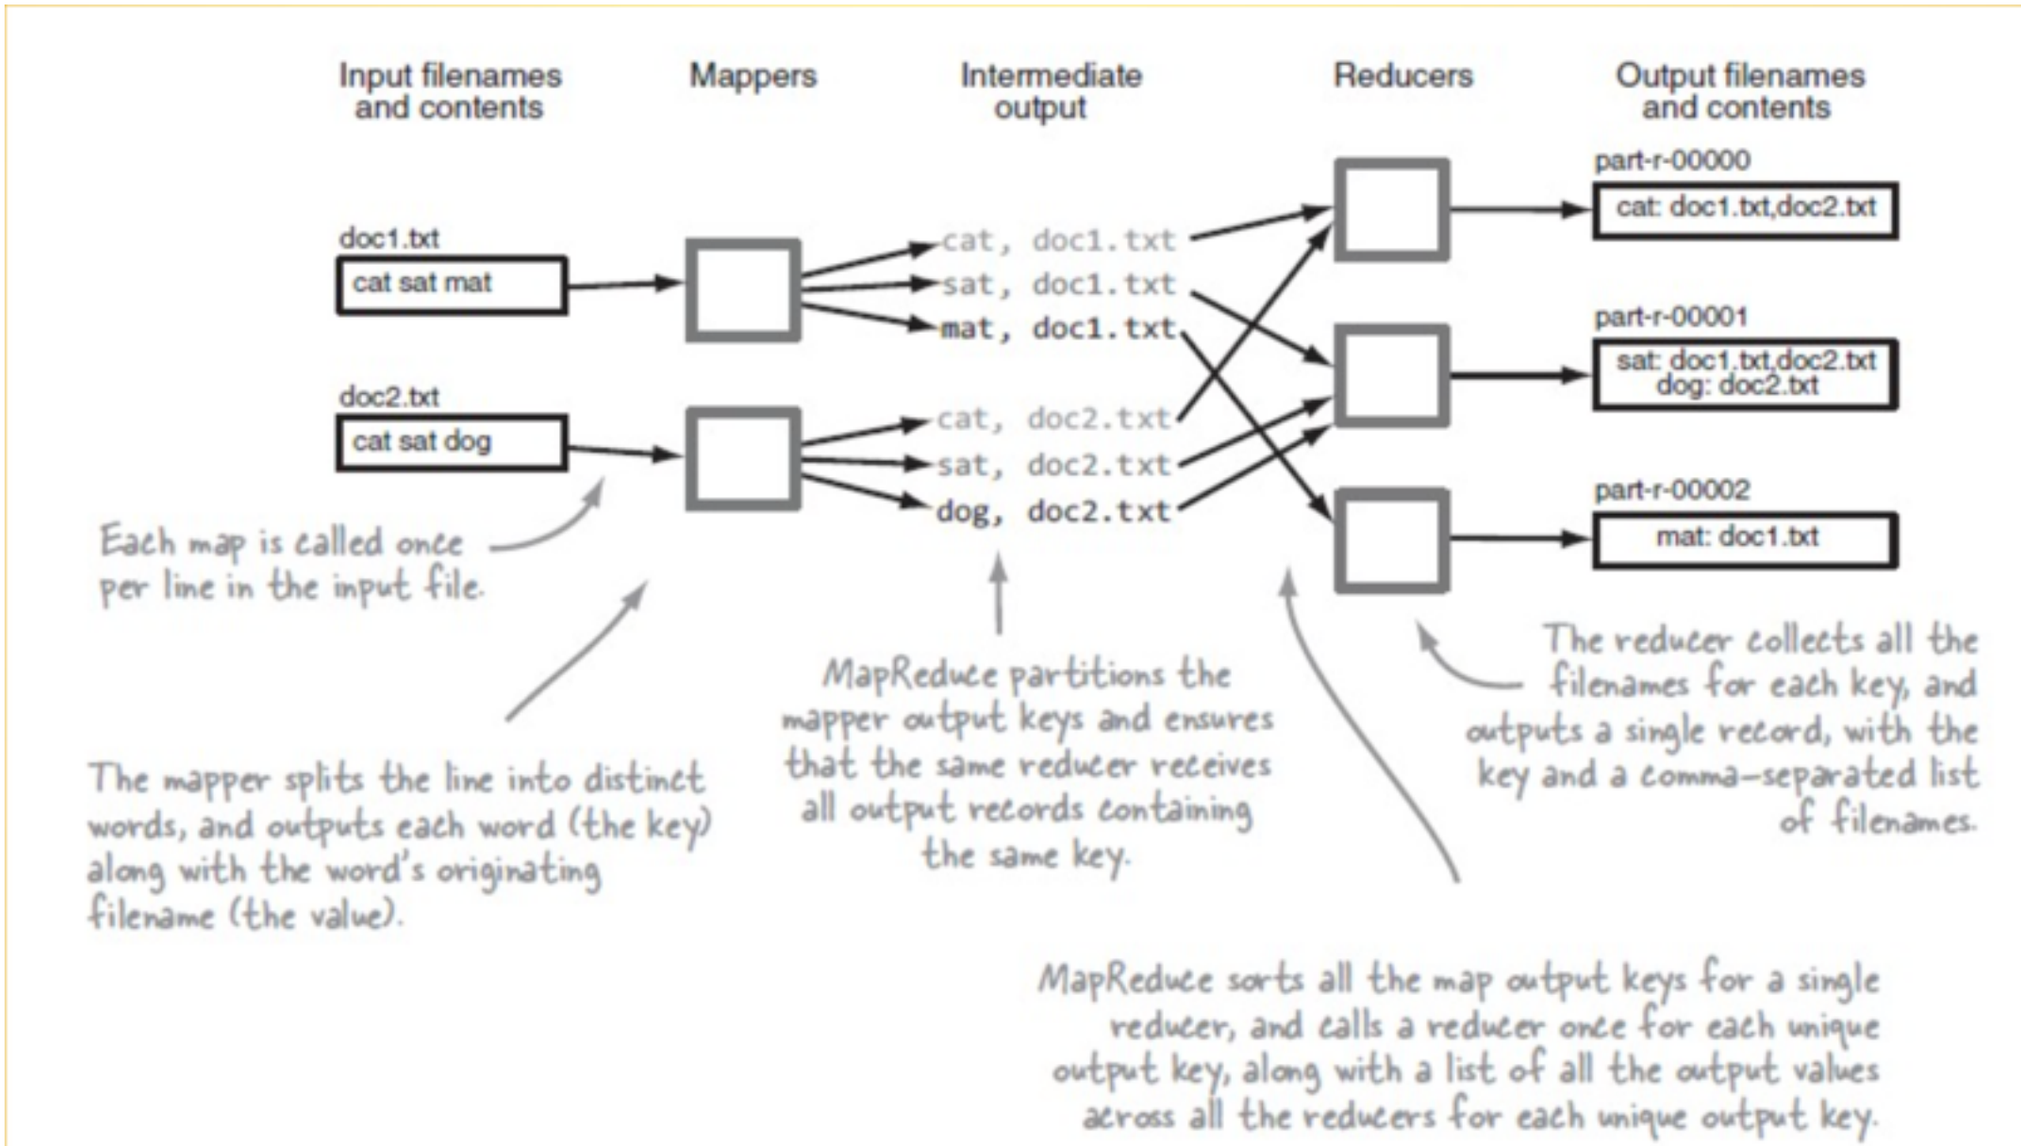

Map Output

The map function takes as input a key/value pair, which represents a logical record from the input data source. In the case of a file, this could be a line, or if the input source is a table in a database, it could be a row.

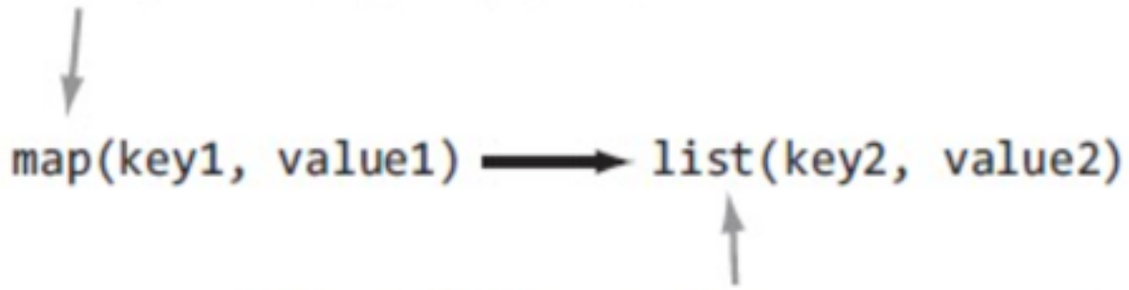

The map function produces zero or more output key/value pairs for that one input pair. For example, if the map function is a filtering map function, it may only produce output if a certain condition is met. Or it could be performing a demultiplexing operation, where a single input key/value yields multiple key/value output pairs.

Reducer Output

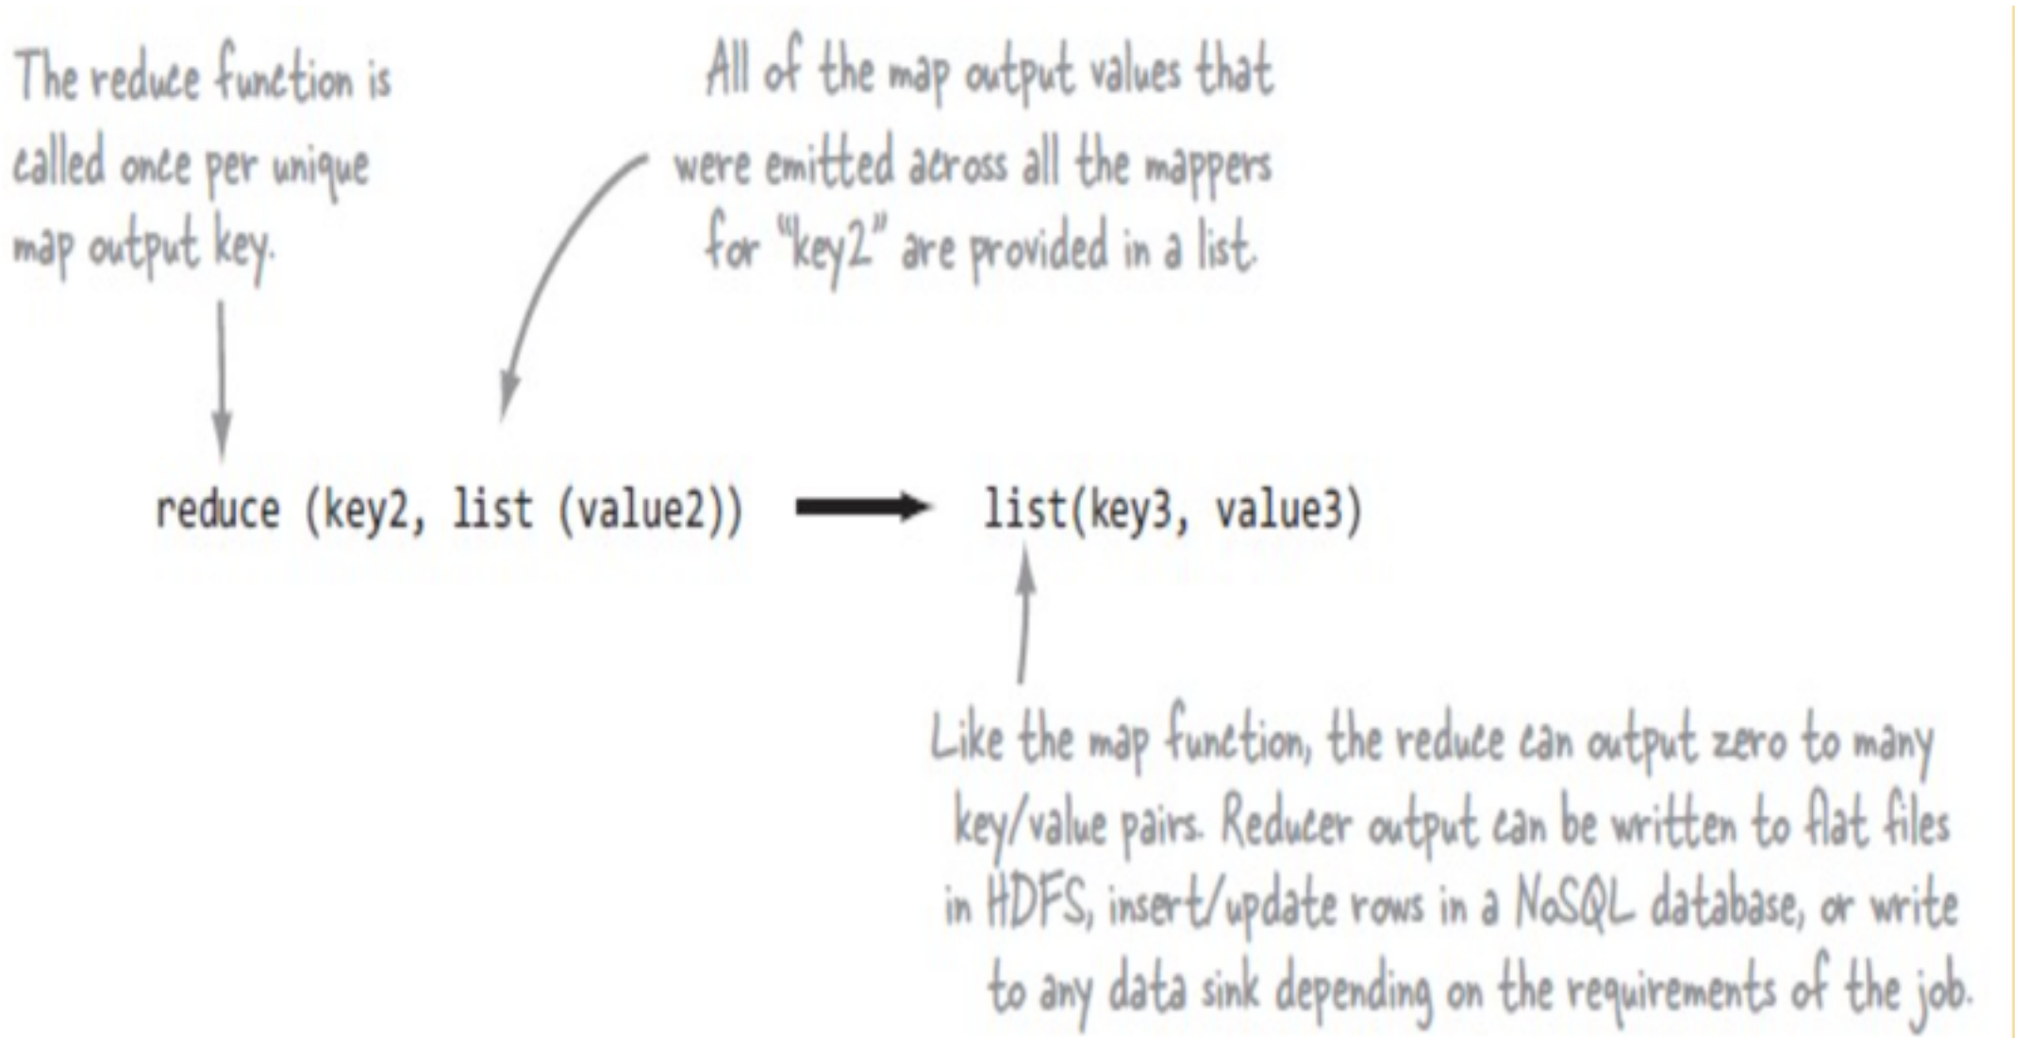

Parallel Execution

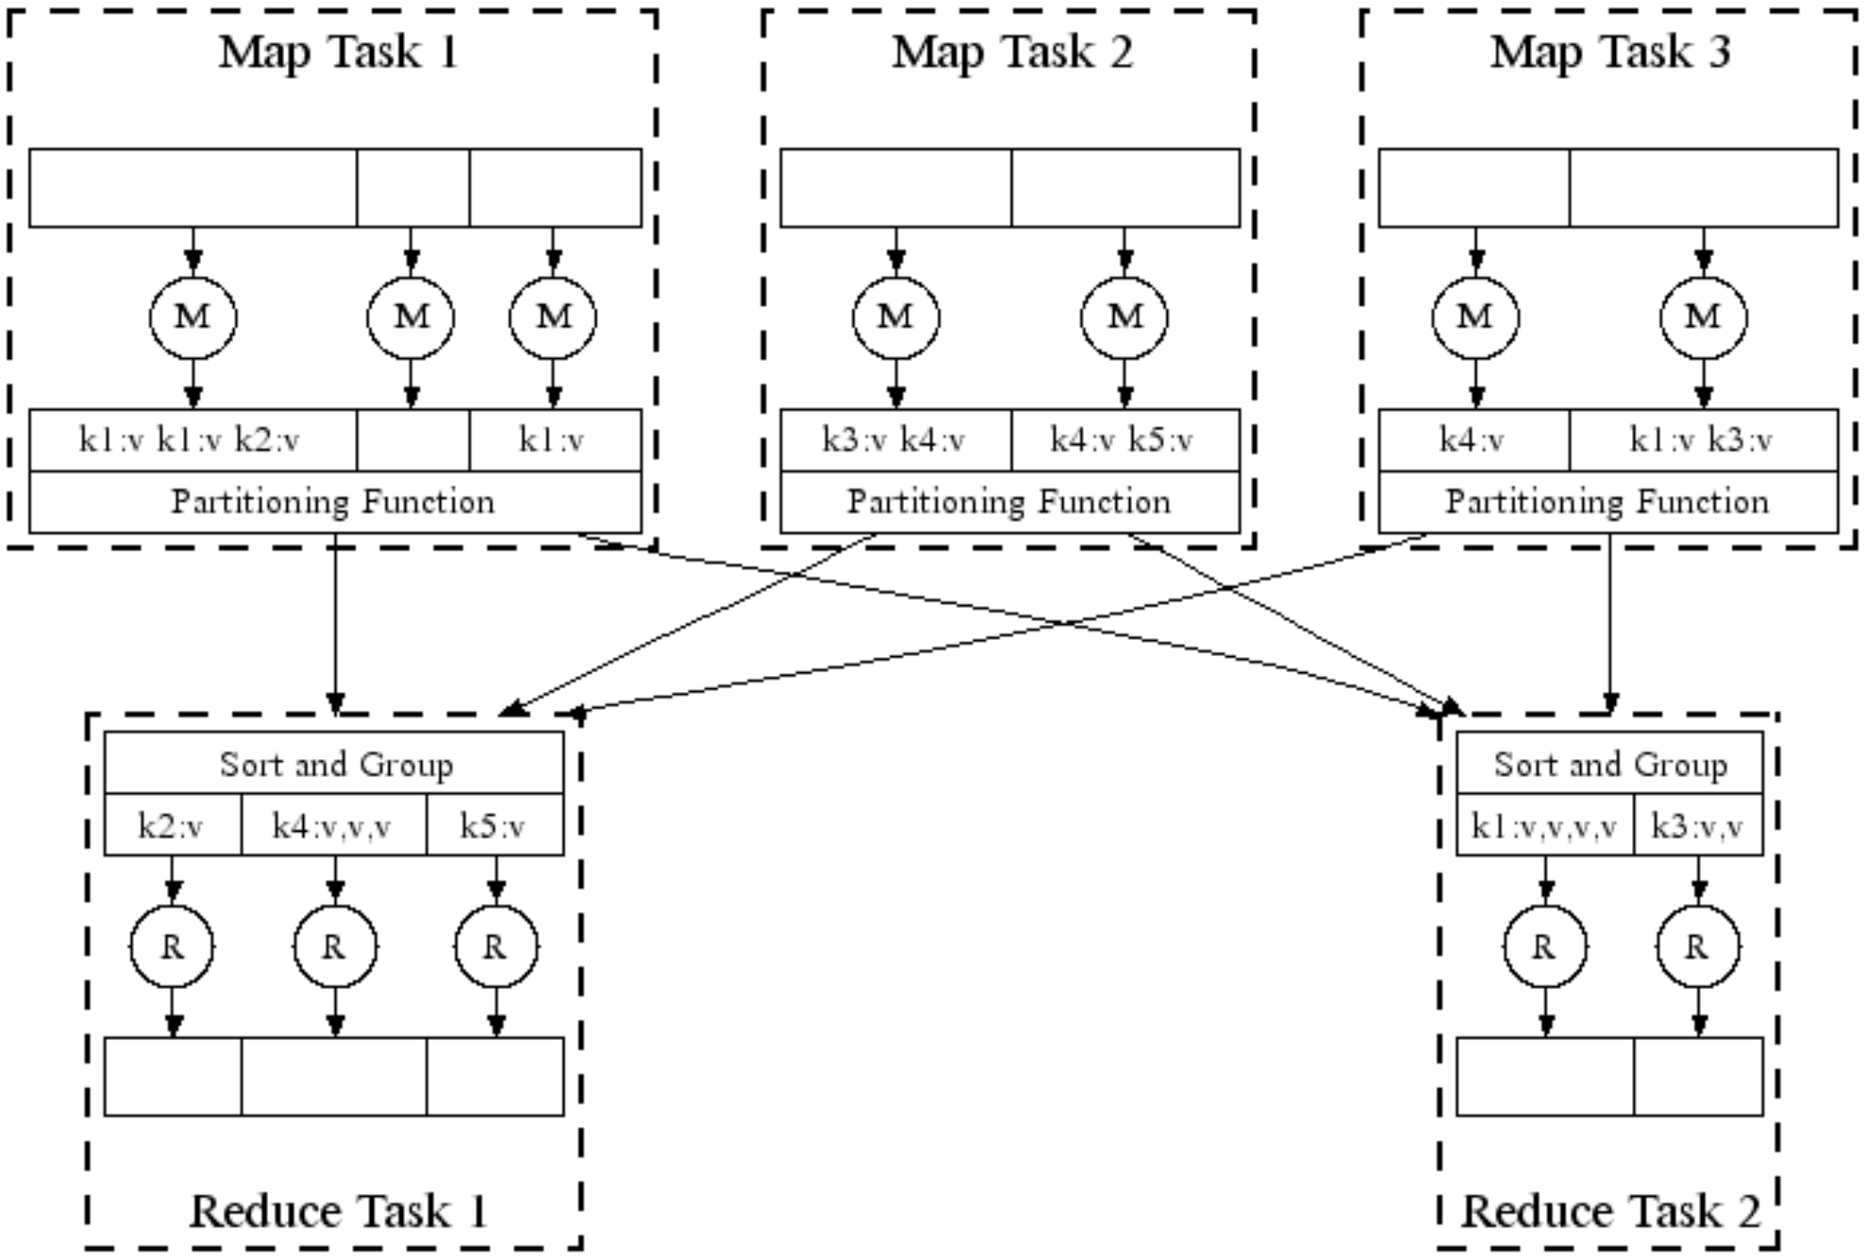

Map Reduce Execution

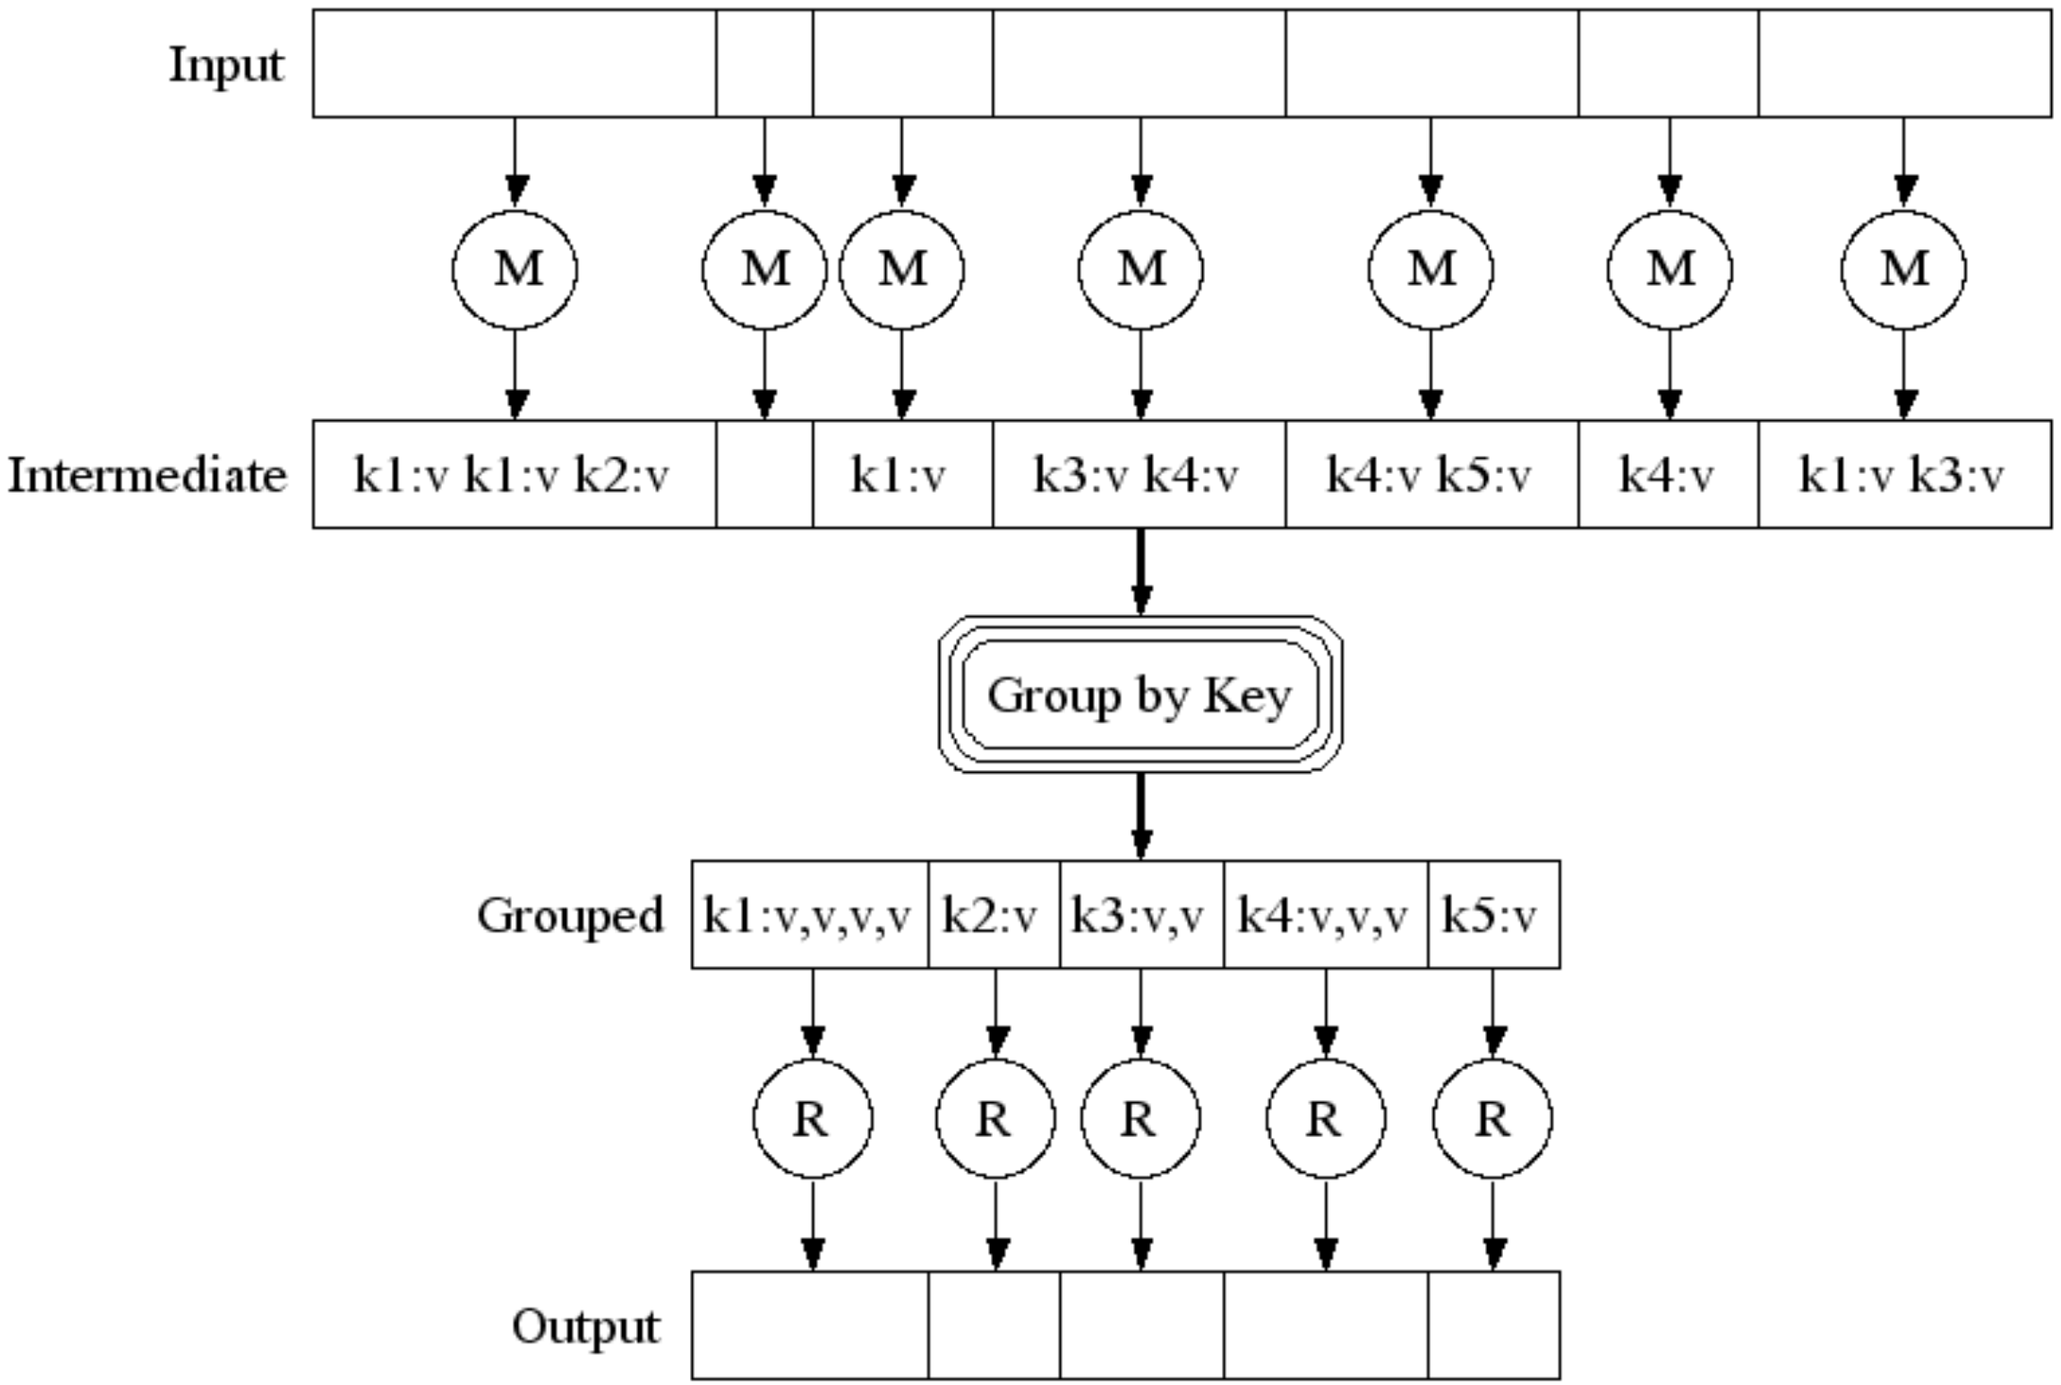

Map Reduce Limitations

- Difficult to Program an algo in Native Map Reduce
- Not good for low latency processing as:
    - Lots of Disk I\O
    - Processing in Stages: Reducer cant start until Mapper is complete
    - Not fit for small batch.
    - Many categories of algos not supported or are very difficult to implement: For e.g., the ones which need iterations
    - The problem of "***stragglers***" (slow workers)
        - Other jobs consuming resources on machine
        - Bad disks with soft errors transfer data very slowly
        - Weird things: processor caches disabled (!!)

→ In short, MR doesn't compose well for large applications

→ We are forced to take "hybrid" approaches many times

→ Therefore, many specialized systems evolved over a period of time as workarounds

What is a DAG - Direct Acyclic Graph Depicting LAZY VAL

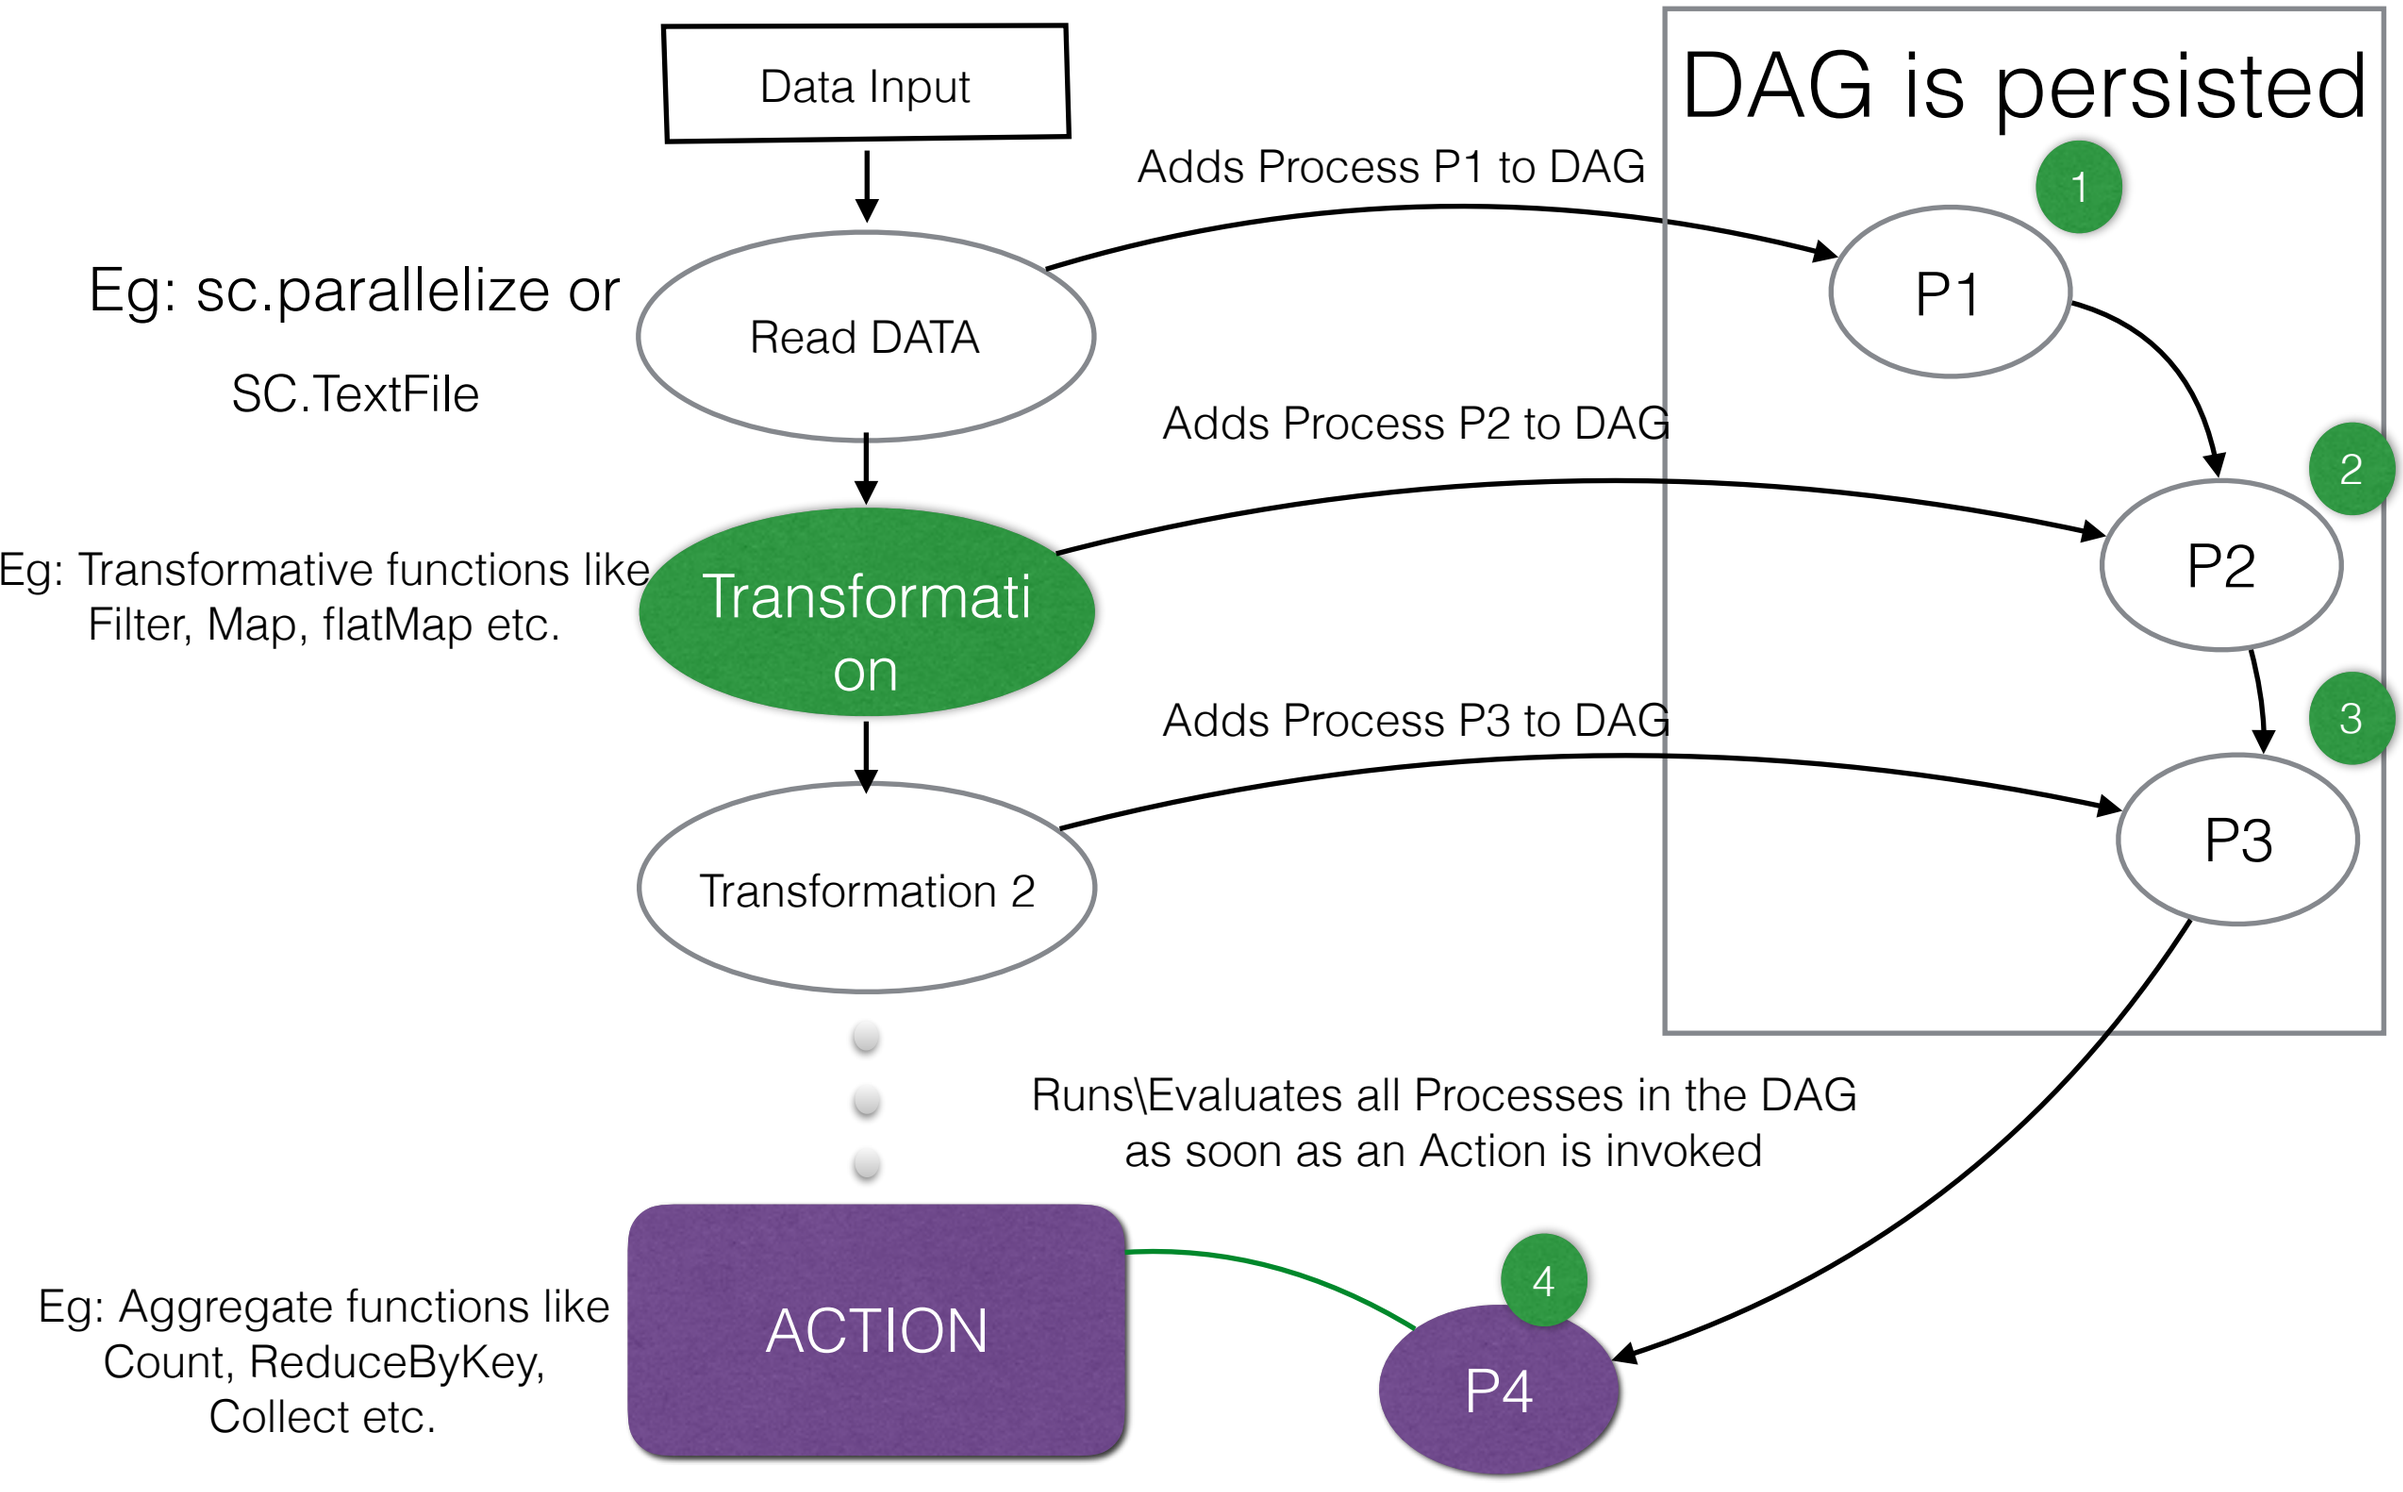

RDD is:

— Smallest unit of Abstracted data set in spark

— it is partitioned so that it can be worked upon in parallel

W1
a = 1
b = 1
c = 1,
c = 1,
a = 1
b = 1

reduceByKey() =>
(a,2),
(b,2),
(c,2

T1 T2 T3

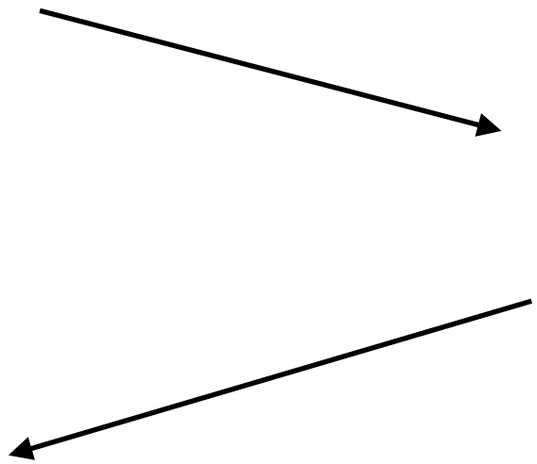

Shuffle

(a,4),(b,4),(c,4)

W2
a = 1
b = 1
c = 1,
c = 1,
a = 1
b = 1

reduceByKey() =>
(a,2),
(b,2),
(c,2

T4 T5 T6

groupByKey()

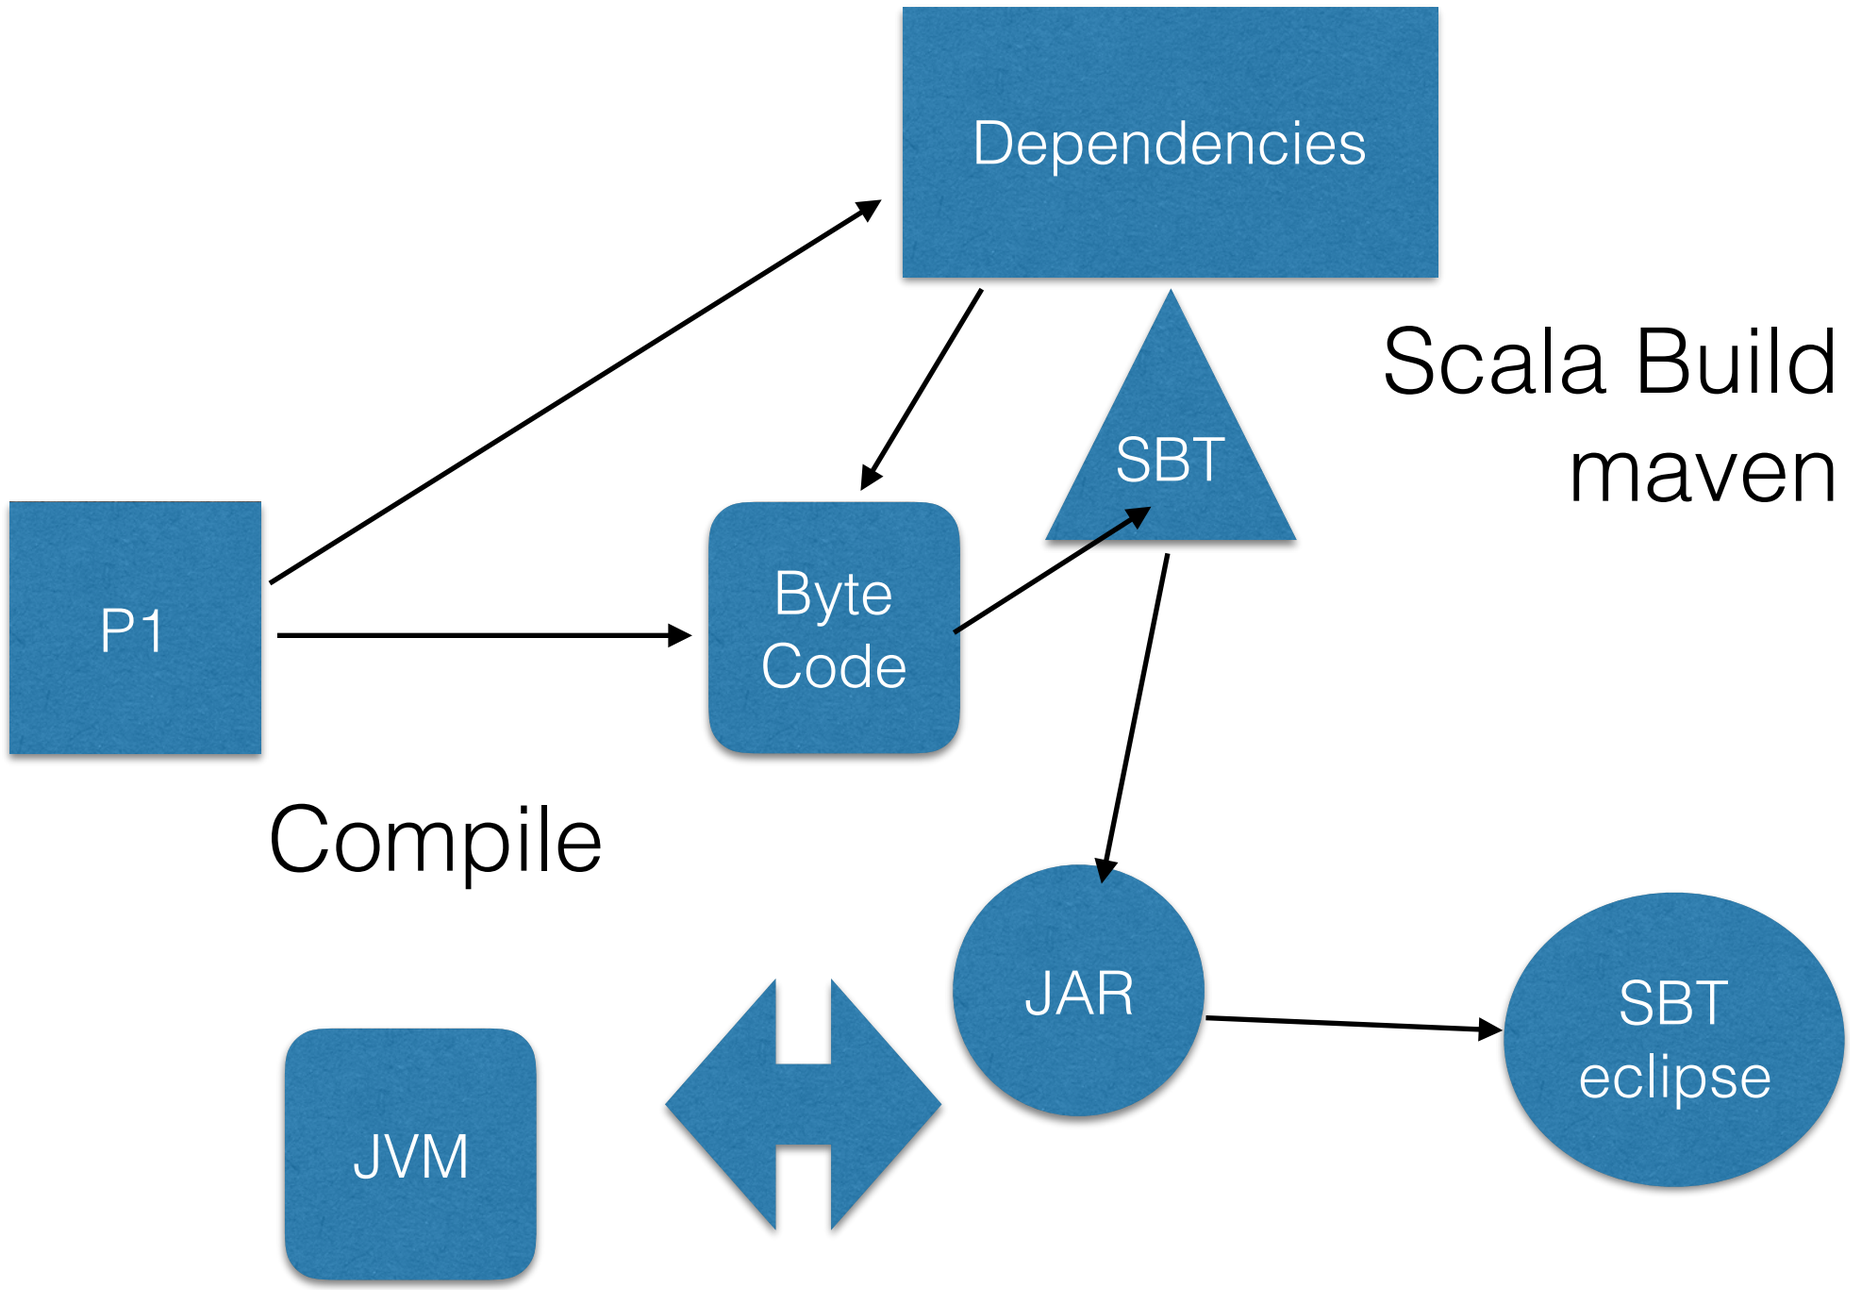
Dependencies
SBT
Scala Build Tool
maven
P1
Byte
Code
Compile
JAR
SBT
eclipse
JVM

build.sbt

1) sbt compile
2) sbt package
3) sbt eclipse

Transformations

Rdd.persist

Filter1

Filter2

Map

HDFS

Join

Collect

Filter1

Filter2

Map

COGROUP